

Assignment_3

Barini Simhadri

10-16-2023

Jump to -
Problem 1
Problem 2
Problem 3
Problem 4

Importing necessary libraries

```
library(rpart)
library(tidyverse)
library(caret)
library(ggplot2)
library(dplyr)
library(rattle)
```

```
data =
read.csv2("C:/Users/bunty/Desktop/funda/week_5/breast_cancer_updated.csv",
header = T, sep = ",")
d = data
```

Problem 1

For this problem, you will perform a straightforward training and evaluation of a decision tree, as well as generate rules by hand. Load the breast_cancer_updated.csv data. These data are visual features computed from samples of breast tissue being evaluated for cancer1. As a preprocessing step, remove the IDNumber column and exclude rows with NA from the dataset.

```
summary(d)
```

```
##      IDNumber      ClumpThickness      UniformCellSize      UniformCellShape
##  Min.   : 61634    Min.   : 1.000    Min.   : 1.000    Min.   : 1.000
## 1st Qu.: 870688    1st Qu.: 2.000    1st Qu.: 1.000    1st Qu.: 1.000
## Median : 1171710    Median : 4.000    Median : 1.000    Median : 1.000
## Mean   : 1071704    Mean   : 4.418    Mean   : 3.134    Mean   : 3.207
## 3rd Qu.: 1238298    3rd Qu.: 6.000    3rd Qu.: 5.000    3rd Qu.: 5.000
## Max.   :13454352    Max.   :10.000    Max.   :10.000    Max.   :10.000
##
##  MarginalAdhesion EpithelialCellSize      BareNuclei      BlandChromatin
```

##	Min.	:	1.000	Min.	:	1.000	Min.	:	1.000	Min.	:	1.000
##	1st Qu.:		1.000	1st Qu.:		2.000	1st Qu.:		1.000	1st Qu.:		2.000
##	Median :		1.000	Median :		2.000	Median :		1.000	Median :		3.000

```
## Mean : 2.807 Mean : 3.216 Mean : 3.545 Mean : 3.438
## 3rd Qu.: 4.000 3rd Qu.: 4.000 3rd Qu.: 6.000 3rd Qu.: 5.000
## Max. :10.000 Max. :10.000 Max. :10.000 Max. :10.000
##
## NormalNucleoli Mitoses Class
## Min. : 1.000 Min. : 1.000 Length:699
## 1st Qu.: 1.000 1st Qu.: 1.000 Class :character
## Median : 1.000 Median : 1.000 Mode :character
## Mean : 2.867 Mean : 1.589
## 3rd Qu.: 4.000 3rd Qu.: 1.000
## Max. :10.000 Max. :10.000
##
```

Removing IDNumber column

```
d = select(d, -c(IDNumber))
```


dropping rows with all NA's and it shows 0 after dropping all.

```
d = na.omit(d)
d[rowSums(is.na(d)) > 0, ]
```

```
## [1] ClumpThickness UniformCellSize UniformCellShape MarginalAdhesion
## [5] EpithelialCellSize BareNuclei BlandChromatin NormalNucleoli
## [9] Mitoses Class
## <0 rows> (or 0-length row.names)
```

a. Apply decision tree learning (use rpart) to the data to predict breast cancer malignancy (Class) and report the accuracy using 10-fold cross validation.

```
ggplot(d,aes(x=Class)) + geom_bar()
```



```
#evaluation method
train_control = trainControl(method = "cv", number = 10)
# Fit the model
tree1 <- train(Class ~., data = d, method = "rpart", trControl = train_control)
# Evaluate fit
tree1
```

```
## CART
##
## 683 samples
## 9 predictor
## 2 classes: 'benign', 'malignant'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 614, 615, 615, 615, 615, 614, ...
## Resampling results across tuning parameters:
##
##   cp          Accuracy    Kappa
## 0.02510460  0.9370205  0.8621292
## 0.05439331  0.9209292  0.8297487
## 0.79079498  0.8287084  0.5567141
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was cp = 0.0251046.
```

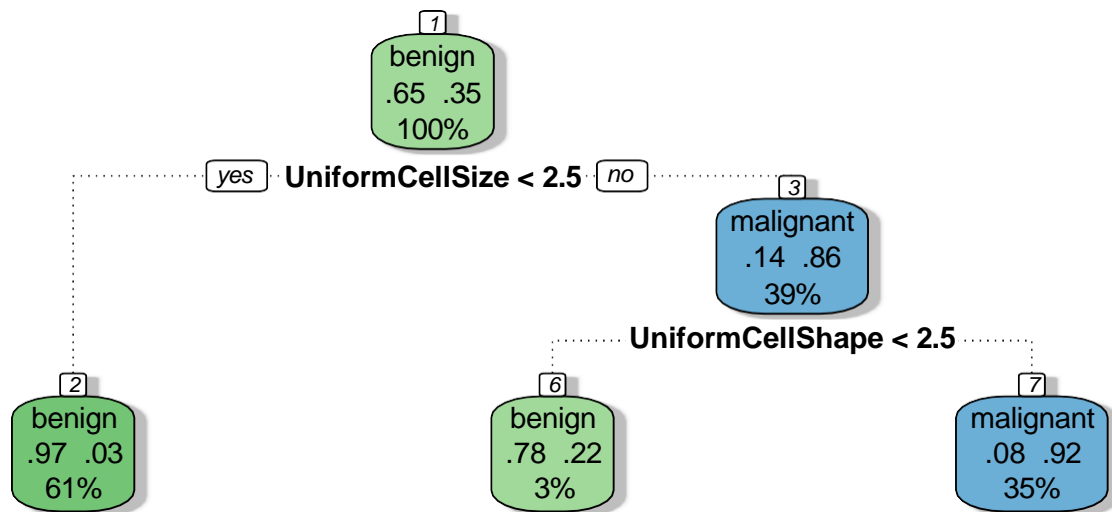
```
pred_tree <- predict(tree1, d)

# Generate confusion matrix for the test set
confusionMatrix(as.factor(d$Class), pred_tree)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  benign malignant
##   benign      424         20
##   malignant    17         222
##
##           Accuracy : 0.9458
##           95% CI : (0.9261, 0.9616)
##   No Information Rate : 0.6457 ##
##   P-Value [Acc > NIR] : <2e-16 ##
##           Kappa : 0.8813
##
##   McNemar's Test P-Value : 0.7423
##
##           Sensitivity : 0.9615
##           Specificity : 0.9174
##           Pos Pred Value : 0.9550
##           Neg Pred Value : 0.9289
##           Prevalence : 0.6457
##           Detection Rate : 0.6208
##   Detection Prevalence : 0.6501 ##
##           Balanced Accuracy : 0.9394
##
##           'Positive' Class : benign
##
```

b. Generate a visualization of the decision tree.

```
fancyRpartPlot(tree1$finalModel, caption = "")
```



generated result using if-then and added the result in separate column.

```

n=1
for(n in 1:nrow(d)){
  x <- d$UniformCellSize[n]
  y <- d$UniformCellShape[n]
  if(x & y >= 2.5){
    result="malignant"
  } else if(x >= 2.5 & y < 2.5){
    result="Bening"
  } else {
    result="Bening"
  }
  d[n,"ifthenresult"]<- result
  n=n+1
}

```

head(d)

##	ClumpThickness	UniformCellSize	UniformCellShape	MarginalAdhesion
## 1	5	1	1	1
## 2	5	4	4	5
## 3	3	1	1	1
## 4	6	8	8	1
## 5	4	1	1	3
## 6	8	10	10	8

```
##      EpithelialCellSize BareNuclei BlandChromatin NormalNucleoli Mitoses      Class
## 1              2          1          3          1          1      benign
## 2              7         10          3          2          1      benign
## 3              2          2          3          1          1      benign
## 4              3          4          3          7          1      benign
## 5              2          1          3          1          1      benign
## 6              7         10          9          7          1 malignant
##      ifthenresult
## 1      Bening
## 2      malignant
## 3      Bening
## 4      malignant
## 5      Bening
## 6      malignant
```

Problem 2

In this problem you will generate decision trees with a set of parameters. You will be using the storms data, a subset of the NOAA Atlantic hurricane database2 , which includes the positions and attributes of 198 tropical storms (potential hurricanes), measured every six hours during the lifetime of a storm. It is part of the dplyr library, so load the library and you will be able to access it. As a preprocessing step, view the data and make sure the target variable (category) is converted to a factor (as opposed to character string).

```
head(storms)
```

```
## # A tibble: 6 x 13
##   name   year month   day   hour   lat   long status   category   wind   pressure
##   <chr> <dbl> <dbl> <int> <dbl> <dbl> <dbl> <fct>      <dbl> <int>    <int>
## 1 Amy   1975     6    27     0  27.5 -79   tropical de~      NA    25    1013
## 2 Amy   1975     6    27     6  28.5 -79   tropical de~      NA    25    1013
## 3 Amy   1975     6    27    12  29.5 -79   tropical de~      NA    25    1013
## 4 Amy   1975     6    27    18  30.5 -79   tropical de~      NA    25    1013
## 5 Amy   1975     6    28     0  31.5 -78.8 tropical de~      NA    25    1012
## 6 Amy   1975     6    28     6  32.4 -78.7 tropical de~      NA    25    1012
## # i 2 more variables: tropicalstorm_force_diameter <int>,
## #   hurricane_force_diameter <int>
```

```
storm_data <- storms
```

```
summary(storm_data)
```

```
##           name           year           month           day
## Length:19066      Min.   :1975      Min.   : 1.000      Min.   : 1.00
## Class :character  1st Qu.:1993      1st Qu.: 8.000      1st Qu.: 8.00
## Mode :character   Median :2004      Median : 9.000      Median :16.00
##                  Mean   :2002      Mean   : 8.699      Mean   :15.78
##                  3rd Qu.:2012      3rd Qu.: 9.000      3rd Qu.:24.00
##                  Max.   :2021      Max.   :12.000      Max.   :31.00
##
##           hour           lat           long           status
## Min.   : 0.000      Min.   : 7.00      Min.   :-109.30      tropical storm      6684
```

```
## 1st Qu.: 5.000    1st Qu.:18.40    1st Qu.: -78.70    hurricane          4684
## Median :12.000    Median :26.60    Median : -62.25    tropical depression:3525
## Mean   : 9.094    Mean   :26.99    Mean   : -61.52    extratropical      2068
## 3rd Qu.:18.000    3rd Qu.:33.70    3rd Qu.: -45.60    other low          1405
## Max.   :23.000    Max.   :70.70    Max.   : 13.50    subtropical storm : 292
##                                     (Other)           : 408
##      category      wind      pressure    tropicalstorm_force_diameter
## Min.   :1.000    Min.   : 10.00    Min.   : 882.0    Min.   : 0.0
## 1st Qu.:1.000    1st Qu.: 30.00    1st Qu.: 987.0    1st Qu.: 0.0
## Median :1.000    Median : 45.00    Median :1000.0    Median : 110.0
## Mean   :1.898    Mean   : 50.02    Mean   : 993.6    Mean   : 146.3
## 3rd Qu.:3.000    3rd Qu.: 65.00    3rd Qu.:1007.0    3rd Qu.: 220.0
## Max.   :5.000    Max.   :165.00    Max.   :1024.0    Max.   :1440.0
## NA's    :14382                                     NA's    :9512
## hurricane_force_diameter
## Min.   : 0.00
## 1st Qu.: 0.00
## Median : 0.00
## Mean   : 14.81
## 3rd Qu.: 0.00
## Max.   :300.00
## NA's    :9512
```

```
str(storm_data)
```

```
## tibble [19,066 × 13] (S3: tbl_df/tbl/data.frame)
## $ name      : chr [1:19066] "Amy" "Amy" "Amy" "Amy" ...
## $ year      : num [1:19066] 1975 1975 1975 1975 1975 ...
## $ month     : num [1:19066] 6 6 6 6 6 6 6 6 6 6 ...
## $ day       : int [1:19066] 27 27 27 27 28 28 28 28 29 29 ...
## $ hour      : num [1:19066] 0 6 12 18 0 6 12 18 0 6 ...
## $ lat       : num [1:19066] 27.5 28.5 29.5 30.5 31.5 32.4 33.3 34 34.4 34 ...
## $ long      : num [1:19066] -79 -79 -79 -79 -78.8 -78.7 -78 -77 -75.8 -74.8 ...
## $ status    : Factor w/ 9 levels "disturbance",...: 7 7 7 7 7 7 7 7 8 8 ...
## $ category  : num [1:19066] NA NA NA NA NA NA NA NA NA NA ...
## $ wind      : int [1:19066] 25 25 25 25 25 25 25 30 35 40 ...
## $ pressure  : int [1:19066] 1013 1013 1013 1013 1012 1012 1011 1006 1004 1002 ...
## $ tropicalstorm_force_diameter: int [1:19066] NA NA NA NA NA NA NA NA NA NA ...
## $ hurricane_force_diameter    : int [1:19066] NA NA NA NA NA NA NA NA NA NA ...
```

```
storm_data$category = as.factor(storm_data$category)
str(storm_data$category)
```

```
## Factor w/ 5 levels "1","2","3","4",...: NA NA NA NA NA NA NA NA NA NA ...
```

```
storm_data = na.omit(storm_data)
```

a. Build a decision tree using the following hyperparameters, maxdepth=2, minsplit=5 and minbucket=3. Be careful to use the right method of training so that you are not automatically tuning the cp parameter, but you are controlling the aforementioned parameters specifically. Use cross validation to report your accuracy score. These parameters will result in a relatively small tree.

Set hyper parameters to controls minsplit, maxdepth, and minbucket


```
hypers = rpart.control(minsplit = 5, maxdepth = 2, minbucket = 3)
```

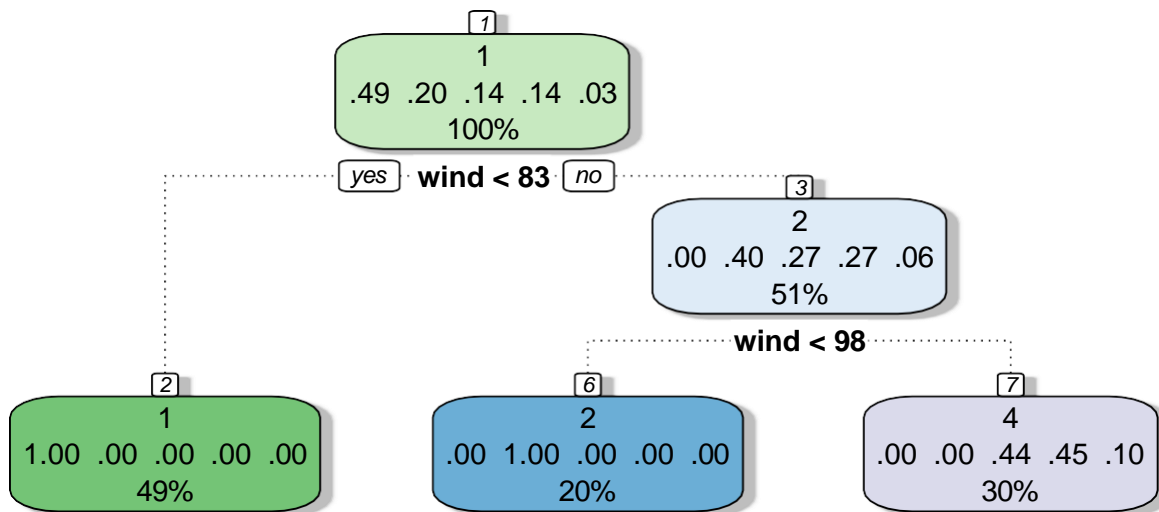
We can see the performance of our model.

```
#fit the model
tree_storm <- train(category ~., data = storm_data, control = hypers, trControl =
  train_control, method = "rpart1SE")

#evaluate
tree_storm
```

```
## CART
##
## 2051 samples
## 12 predictor
## 5 classes: '1', '2', '3', '4', '5'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 1846, 1846, 1847, 1845, 1846, 1846, ...
## Resampling results:
##
## Accuracy Kappa
## 0.8337486 0.7530905
```

```
fancyRpartPlot(tree_storm$finalModel, caption = "")
```



```
pred_tree_storm = predict(tree_storm,storm_data)
confusionMatrix(storm_data$category,pred_tree_storm)
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction      1      2      3      4      5
```

```
##           1 1013      0      0      0      0
```

```
##           2   0  414      0      0      0
```

```
##           3   0   0      0  277      0
```

```
##           4   0   0      0  283      0
```

```
##           5   0   0      0   64      0
```

```
##
```

```
## Overall Statistics
```

```
##
```

```
##           Accuracy : 0.8337
```

```
##           95% CI : (0.8169, 0.8496)
```

```
##           No Information Rate : 0.4939
```

```
##           P-Value [Acc > NIR] : < 2.2e-16
```

```
##
```

```
##           Kappa : 0.7531
```

```
##
```

```
##           McNemar's Test P-Value : NA
```

```
##
```

```
## Statistics by Class:
```

```
##
##
## Class: 1 Class: 2 Class: 3 Class: 4 Class: 5
## Sensitivity      1.0000 1.0000 NA 0.4535 NA
## Specificity      1.0000 1.0000 0.8649 1.0000 0.9688
## Pos Pred Value   1.0000 1.0000 NA 1.0000 NA
## Neg Pred Value    1.0000 1.0000 NA 0.8071 NA
## Prevalence       0.4939 0.2019 0.0000 0.3042 0.0000
## Detection Rate    0.4939 0.2019 0.0000 0.1380 0.0000
## Detection Prevalence 0.4939 0.2019 0.1351 0.1380 0.0312
## Balanced Accuracy 1.0000 1.0000 NA 0.7268 NA
```

b. To see how this performed with respect to the individual classes, we could use a confusion matrix. We also want to see if that aspect of performance is different on the train versus the test set. Create a train/test partition. Train on the training set. By making predictions with that model on the train set and on the test set separately, use the outputs to create two separate confusion matrices, one for each partition. Remember, we are testing if the model built with the training data performs differently on data used to train it (train set) as opposed to new data (test set). Compare the confusion matrices and report which classes it has problem classifying. Do you think that both are performing similarly and what does that suggest about overfitting for the model?

partitioning of data.

```
index = createDataPartition(y=storm_data$category, p=0.7, list=FALSE)
```

```
train_set = storm_data[index,]
test_set = storm_data[-index,]
```

```
#fit the model
tree3 <- train(category ~., data = train_set, control = hypers, trControl =
  train_control, method = "rpart1SE")
```

```
tree3
```

```
## CART
##
## 1438 samples
## 12 predictor
## 5 classes: '1', '2', '3', '4', '5'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 1294, 1295, 1295, 1294, 1295, 1293, ...
## Resampling results:
##
## Accuracy Kappa
## 0.8338124 0.7532243
```

```
pred_tree_train = predict(tree3,train_set)
pred_tree_test = predict(tree3,test_set)
```

```
confusionMatrix(train_set$category,pred_tree_train)
```

```
## Confusion Matrix and Statistics
```

```
##
##           Reference
## Prediction  1    2    3    4    5
##           1 710    0    0    0    0
##           2   0 290    0    0    0
##           3   0   0    0 194    0
##           4   0   0    0 199    0
##           5   0   0    0  45    0
##
```

```
## Overall Statistics
```

```
##
##           Accuracy : 0.8338
##           95% CI : (0.8135, 0.8527)
##           No Information Rate : 0.4937
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.7532
##
```

```
## McNemar's Test P-Value : NA
```

```
##
## Statistics by Class:
```

```
##
##           Class: 1 Class: 2 Class: 3 Class: 4 Class: 5
## Sensitivity      1.0000   1.0000      NA   0.4543      NA
## Specificity      1.0000   1.0000   0.8651   1.0000   0.96871
## Pos Pred Value   1.0000   1.0000      NA   1.0000      NA
## Neg Pred Value    1.0000   1.0000      NA   0.8071      NA
## Prevalence       0.4937   0.2017   0.0000   0.3046   0.00000
## Detection Rate    0.4937   0.2017   0.0000   0.1384   0.00000
## Detection Prevalence 0.4937   0.2017   0.1349   0.1384   0.03129
## Balanced Accuracy 1.0000   1.0000      NA   0.7272      NA
```

```
confusionMatrix(test_set$category,pred_tree_test)
```

```
## Confusion Matrix and Statistics
```

```
##
##           Reference
## Prediction  1    2    3    4    5
##           1 303    0    0    0    0
##           2   0 124    0    0    0
##           3   0   0    0  83    0
##           4   0   0    0  84    0
##           5   0   0    0  19    0
##
```

```
## Overall Statistics
```

```
##
##           Accuracy : 0.8336
##           95% CI : (0.8017, 0.8622)
##           No Information Rate : 0.4943
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.7528
##
```

```
## McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: 1 Class: 2 Class: 3 Class: 4 Class: 5
## Sensitivity      1.0000  1.0000      NA   0.4516      NA
## Specificity      1.0000  1.0000  0.8646  1.0000  0.969
## Pos Pred Value   1.0000  1.0000      NA   1.0000      NA
## Neg Pred Value    1.0000  1.0000      NA   0.8072      NA
## Prevalence        0.4943  0.2023  0.0000  0.3034  0.000
## Detection Rate    0.4943  0.2023  0.0000  0.1370  0.000
## Detection Prevalence 0.4943  0.2023  0.1354  0.1370  0.031
## Balanced Accuracy 1.0000  1.0000      NA   0.7258      NA
```

From above two confusion matrices , we can see the accuracy is almost same for predicting both on train and set set which indicates there is no overfitting. The model misclassify to predict some intances of class 3 and 5.

Problem 3

This is will be an extension of Problem 2, using the same data and class. Here you will build many decision trees, manually tuning the parameters to gain intuition about the tradeoffs and how these tree parameters affect the complexity and quality of the model. The goal is to find the best tree model, which means it should be accurate but not too complex that the model overfits the training data. We will achieve this by using multiple sets of parameters and creating a graph of accuracy versus complexity for the training and the test sets (refer to the tutorial). This problem may require a significant amount of effort because you will need to train a substantial number of trees (at least 10).

a. Partition your data into 80% for training and 20% for the test data set

```
index = createDataPartition(y=storm_data$category, p=0.8, list=FALSE)
train_set = storm_data[index,]
test_set = storm_data[-index,]
```

```
str(test_set)
```

```
## tibble [407 x 13] (S3: tbl_df/tbl/data.frame)
## $ name      : chr [1:407] "Alex" "Charley" "Danielle" "Danielle" ...
## $ year      : num [1:407] 2004 2004 2004 2004 2004 ...
## $ month     : num [1:407] 8 8 8 8 8 8 8 8 8 8 ...
## $ day       : int [1:407] 5 14 15 15 17 18 27 28 30 30 ...
## $ hour      : num [1:407] 0 12 6 18 6 0 6 12 6 18 ...
## $ lat       : num [1:407] 38.5 32.3 14.1 15.2 21.7 25.9 14.2 17.2 19 19.4 ...
## $ long      : num [1:407] -66 -79.7 -30.8 -33.5 -39.6 -40.6 -47.8 -51.6 -56.8 -59
## $ status    : Factor w/ 9 levels "disturbance",...: 3 3 3 3 3 3 3 3 3 ...
## $ category  : Factor w/ 5 levels "1","2","3","4",...: 3 1 1 2 2 1 1 3 3 3 ...
## $ wind      : int [1:407] 105 65 75 85 90 75 75 105 100 110 ...
## $ pressure  : int [1:407] 957 988 981 975 970 981 980 958 958 948 ...
## $ tropicalstorm_force_diameter: int [1:407] 225 140 195 195 150 150 135 180 260 280 ...
## $ hurricane_force_diameter   : int [1:407] 70 40 45 50 40 40 30 50 70 90 ...
## - attr(*, "na.action")= 'omit' Named int [1:17015] 1 2 3 4 5 6 7 8 9 10 ...
## ..- attr(*, "names")= chr [1:17015] "1" "2" "3" "4" ...
```

b. Train at least 10 trees using different sets of parameters, through you made need more. Create the graph described above such that you can identify the inflection point where the tree is overfitting and pick a high-quality decision tree. Your strategy should be to make at least one very simple model and at least one very complex model and work towards the center by changing different parameters. Generate a table that contains all of the parameters (maxdepth, minsplit, minbucket, etc) used along with the number of nodes created, and the training and testing set accuracy values. The number of rows will be equal to the number of sets of parameters used. You will use the data in the table to generate the graph. The final results to be reported for this problem are the table and graph.

Implementing this question as I learned from the tutorial 6.

Tree 1

```
# Initialize cross validation
train_control = trainControl(method = "cv", number = 10)

# tree 1
hypers = rpart.control(minsplit = 2, maxdepth = 1, minbucket = 2)
tree1 <- train(category ~., data = train_set, control = hypers,
               trControl = train_control, method = "rpart1SE")

# Training set
pred_tree <- predict(tree1, train_set)
cfm_train <- confusionMatrix(train_set$category, pred_tree)

# Test set
pred_tree <- predict(tree1, test_set)
cfm_test <- confusionMatrix(test_set$category, pred_tree)

a_train <- cfm_train$overall[1]
a_test <- cfm_test$overall[1]
nodes <- nrow(tree1$finalModel$frame)

comp_tbl <- data.frame("Nodes" = nodes, "TrainAccuracy" = a_train, "TestAccuracy" = a_test,
                      "MaxDepth" = 1, "Minsplit" = 2, "Minbucket" = 2)
```

Tree 2

```
hypers = rpart.control(minsplit = 5, maxdepth = 2, minbucket = 5)
tree2 <- train(category ~., data = train_set, control = hypers,
               trControl = train_control, method = "rpart1SE")

# Training set
pred_tree <- predict(tree2, train_set)
cfm_train <- confusionMatrix(train_set$category, pred_tree)

# Test set
pred_tree <- predict(tree2, test_set)
cfm_test <- confusionMatrix(test_set$category, pred_tree)

a_train <- cfm_train$overall[1]
a_test <- cfm_test$overall[1]
nodes <- nrow(tree2$finalModel$frame)
```

```
comp_tbl <- rbind(comp_tbl,c(nodes,a_train,a_test,2,5,5))
```

Tree 3

```
hypers = rpart.control(minsplit = 50, maxdepth = 3, minbucket = 50)
tree3 <- train(category ~., data = train_set, control = hypers,
               trControl = train_control, method = "rpart1SE")

# Training set
pred_tree <- predict(tree3, train_set)
cfm_train <- confusionMatrix(train_set$category, pred_tree)

# Test set
pred_tree <- predict(tree3, test_set)
cfm_test <- confusionMatrix(test_set$category, pred_tree)

a_train <- cfm_train$overall[1]
a_test <- cfm_test$overall[1]
nodes <- nrow(tree3$finalModel$frame)

comp_tbl <- rbind(comp_tbl,c(nodes,a_train,a_test,3,50,50))
```

Tree 4

```
hypers = rpart.control(minsplit = 100, maxdepth = 4, minbucket = 100)
tree4 <- train(category ~., data = train_set, control = hypers,
               trControl = train_control, method = "rpart1SE")

# Training set
pred_tree <- predict(tree4, train_set)
cfm_train <- confusionMatrix(train_set$category, pred_tree)

# Test set
pred_tree <- predict(tree4, test_set)
cfm_test <- confusionMatrix(test_set$category, pred_tree)

a_train <- cfm_train$overall[1]
a_test <- cfm_test$overall[1]
nodes <- nrow(tree4$finalModel$frame)

comp_tbl <- rbind(comp_tbl,c(nodes,a_train,a_test,4,100,100))
```

Tree 5

```
hypers = rpart.control(minsplit = 100, maxdepth = 5, minbucket = 100)
tree5 <- train(category ~., data = train_set, control = hypers,
               trControl = train_control, method = "rpart1SE")
```

```

# Training set
pred_tree <- predict(tree5, train_set)
cfm_train <- confusionMatrix(train_set$category, pred_tree)

# Test set
pred_tree <- predict(tree5, test_set)
cfm_test <- confusionMatrix(test_set$category, pred_tree)

a_train <- cfm_train$overall[1]
a_test <- cfm_test$overall[1]
nodes <- nrow(tree5$finalModel$frame)

comp_tbl <- rbind(comp_tbl,c(nodes,a_train,a_test,5,100,100))

```

Tree 6

```

hypers = rpart.control(minsplit = 1000, maxdepth = 4, minbucket = 1000)
tree6 <- train(category ~., data = train_set, control = hypers,
               trControl = train_control, method = "rpart1SE")

# Training set
pred_tree <- predict(tree6, train_set)
cfm_train <- confusionMatrix(train_set$category, pred_tree)

# Test set
pred_tree <- predict(tree6, test_set)
cfm_test <- confusionMatrix(test_set$category, pred_tree)

a_train <- cfm_train$overall[1]
a_test <- cfm_test$overall[1]
nodes <- nrow(tree6$finalModel$frame)

comp_tbl <- rbind(comp_tbl,c(nodes,a_train,a_test,4,1000,1000))

```

Tree 7

```

hypers = rpart.control(minsplit = 2000, maxdepth = 5, minbucket = 2000)
tree7 <- train(category ~., data = train_set, control = hypers,
               trControl = train_control, method = "rpart1SE")

# Training set
pred_tree <- predict(tree7, train_set)
cfm_train <- confusionMatrix(train_set$category, pred_tree)

# Test set
pred_tree <- predict(tree7, test_set)
cfm_test <- confusionMatrix(test_set$category, pred_tree)

a_train <- cfm_train$overall[1]

```



```

a_test <- cfm_test$overall[1]
nodes <- nrow(tree7$finalModel$frame)

comp_tbl <- rbind(comp_tbl,c(nodes,a_train,a_test,5,2000,2000))

```

Tree 8

```

hypers = rpart.control(minsplit = 5000, maxdepth = 8, minbucket = 5000)
tree8 <- train(category ~., data = train_set, control = hypers,
               trControl = train_control, method = "rpart1SE")

# Training set
pred_tree <- predict(tree8, train_set)
cfm_train <- confusionMatrix(train_set$category, pred_tree)

# Test set
pred_tree <- predict(tree8, test_set)
cfm_test <- confusionMatrix(test_set$category, pred_tree)

a_train <- cfm_train$overall[1]
a_test <- cfm_test$overall[1]
nodes <- nrow(tree8$finalModel$frame)

comp_tbl <- rbind(comp_tbl,c(nodes,a_train,a_test,8,5000,5000))

```

Tree 9

```

hypers = rpart.control(minsplit = 10000, maxdepth = 12, minbucket = 10000)
tree9 <- train(category ~., data = train_set, control = hypers,
               trControl = train_control, method = "rpart1SE")

# Training set
pred_tree <- predict(tree9, train_set)
cfm_train <- confusionMatrix(train_set$category, pred_tree)

# Test set
pred_tree <- predict(tree9, test_set)
cfm_test <- confusionMatrix(test_set$category, pred_tree)

a_train <- cfm_train$overall[1]
a_test <- cfm_test$overall[1]
nodes <- nrow(tree9$finalModel$frame)

comp_tbl <- rbind(comp_tbl,c(nodes,a_train,a_test,12,10000,10000))

```

Tree 10

```

hypers = rpart.control(minsplit = 10000, maxdepth = 25, minbucket = 10000)
tree10 <- train(category ~., data = train_set, control = hypers,
                trControl = train_control, method = "rpart1SE")

# Training set
pred_tree <- predict(tree10, train_set)
cfm_train <- confusionMatrix(train_set$category, pred_tree)

# Test set
pred_tree <- predict(tree10, test_set)
cfm_test <- confusionMatrix(test_set$category, pred_tree)

a_train <- cfm_train$overall[1]
a_test <- cfm_test$overall[1]
nodes <- nrow(tree10$finalModel$frame)

comp_tbl <- rbind(comp_tbl, c(nodes, a_train, a_test, 25, 10000, 10000))

comp_tbl


```

##	Nodes	TrainAccuracy	TestAccuracy	MaxDepth	Minsplit	Minbucket
## Accuracy	3	0.6952555	0.6977887	1	2	2
## 2	5	0.8333333	0.8353808	2	5	5
## 3	7	0.9683698	0.9705160	3	50	50
## 4	7	0.9683698	0.9705160	4	100	100
## 5	7	0.9683698	0.9705160	5	100	100
## 6	1	0.4933090	0.4963145	4	1000	1000
## 7	1	0.4933090	0.4963145	5	2000	2000
## 8	1	0.4933090	0.4963145	8	5000	5000
## 9	1	0.4933090	0.4963145	12	10000	10000
## 10	1	0.4933090	0.4963145	25	10000	10000

```

#plot the graph
ggplot(comp_tbl, aes(x=Nodes)) +
  geom_point(aes(y = TrainAccuracy), color = "red") +
  geom_point(aes(y = TestAccuracy), color="blue") +
  ylab("Accuracy")

```



c. Identify the final choice of model, list its parameters and evaluate with a confusion matrix to make sure that it gets balanced performance over classes. Also get a better accuracy estimate for this tree using cross validation.

From the above, we can see the Model Tree3 has the highest accuracy with $\text{maxdepth} = 3$, $\text{minsplit} = 50$ and $\text{minbucket} = 50$.

```
#tree 3
hypers = rpart.control(minsplit = 50, maxdepth = 3, minbucket = 50)
tree3 <- train(category ~., data = train_set, control = hypers, trControl = train_control, method = "rpart")
tree3
```

```
## CART
##
## 1644 samples
## 12 predictor
## 5 classes: '1', '2', '3', '4', '5'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 1479, 1480, 1480, 1479, 1480, 1478, ...
## Resampling results:
##
## Accuracy Kappa
## 0.9683735 0.9530949
```

```
# Test set
pred_tree <- predict(tree3, test_set)
confusionMatrix(test_set$category, pred_tree)
```

```
## Confusion Matrix and Statistics
```

```
##
##           Reference
## Prediction   1    2    3    4    5
##           1 202    0    0    0    0
##           2   0   82    0    0    0
##           3   0    0   55    0    0
##           4   0    0    0   56    0
##           5   0    0    0   12    0
##
```

```
## Overall Statistics
```

```
##
##           Accuracy : 0.9705
##           95% CI : (0.9491, 0.9847)
##           No Information Rate : 0.4963
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9561
##
##           McNemar's Test P-Value : NA
##
```

```
## Statistics by Class:
```

```
##
##           Class: 1 Class: 2 Class: 3 Class: 4 Class: 5
## Sensitivity      1.0000   1.0000   1.0000   0.8235      NA
## Specificity      1.0000   1.0000   1.0000   1.0000   0.97052
## Pos Pred Value   1.0000   1.0000   1.0000   1.0000      NA
## Neg Pred Value   1.0000   1.0000   1.0000   0.9658      NA
## Prevalence       0.4963   0.2015   0.1351   0.1671   0.00000
## Detection Rate   0.4963   0.2015   0.1351   0.1376   0.00000
## Detection Prevalence 0.4963   0.2015   0.1351   0.1376   0.02948
## Balanced Accuracy 1.0000   1.0000   1.0000   0.9118      NA
```

Problem 4

In this problem you will identify the most important independent variables used in a classification model. Use the Bank_Modified.csv data. As a preprocessing step, remove the ID column and make sure to convert the target variable, approval, from a string to a factor.

```
bank = read.csv("C:/Users/bunty/Desktop/funda/week_5/Bank_Modified.csv",
header = T, sep = ",")
```

```
##   X cont1 cont2 cont3 bool1 bool2 cont4 bool3 cont5 cont6 approval credit.score
## 1 1 30.83 0.000 1.25    t    t    1    f   202    0    +    664.60
## 2 2 58.67 4.460 3.04    t    t    6    f    43   560    +    693.88
## 3 3 24.50 0.500 1.50    t    f    0    f   280   824    +    621.82
## 4 4 27.83 1.540 3.75    t    t    5    t   100    3    +    653.97
```

```
## 5 5 20.17 5.625 1.71      t      f      0      f    120      0      +      670.26
## 6 6 32.08 4.000 2.50      t      f      0      t    360      0      +      672.16
##   ages
## 1    58
## 2    54
## 3    62
## 4    51
## 5    58
## 6    37
```

```
bank = select(bank , -c("X"))
```

converting to factor

```
bank$approval = as.factor(bank$approval)
str(bank$approval)
```

```
## Factor w/ 2 levels "-","+": 2 2 2 2 2 2 2 2 2 2 2 ...
```

Removing NA's

```
bank = na.omit(bank)
bank[rowSums(is.na(bank)) > 0, ]
```

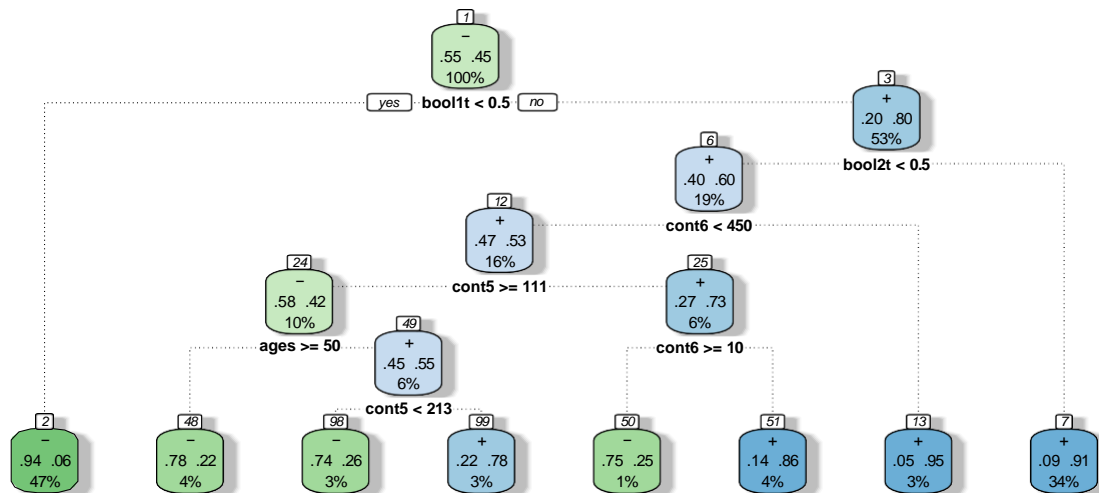
```
## [1] cont1      cont2      cont3      bool1      bool2
## [6] cont4      bool3      cont5      cont6      approval
## [11] credit.score ages
## <0 rows> (or 0-length row.names)
```

a. Build your initial decision tree model with minsplit=10 and maxdepth=20

```
#set hyperparameters
hypers = rpart.control(minsplit = 10, maxdepth = 20)
tree <- train(approval ~., data = bank, control = hypers, trControl =
  train_control, method = "rpart1SE")
tree
```

```
## CART
##
## 666 samples
## 11 predictor
## 2 classes: '-','+'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 600, 599, 600, 599, 599, 600, ...
## Resampling results:
##
## Accuracy   Kappa
## 0.8604025  0.7178918
```

```
fancyRpartPlot(tree$finalModel, caption= "")
```




b. Run variable importance analysis on the model and print the result.

```
var_imp <- varImp(tree, scale= FALSE)
var_imp
```

```
## rpart1SE variable importance
##
##           Overall
## bool1t      179.282
## cont4       97.700
## bool2t      85.622
## ages        72.800
## cont3       64.343
## cont6       26.828
## cont5       25.878
## cont2        9.620
## cont1        5.646
## credit.score 1.504
## bool3t        0.000
```

c. Generate a plot to visualize the variables by importance.

```
plot(var_imp)
```



d . Rebuild your model with the top six variables only, based on the variable relevance analysis. Did this change have an effect on the accuracy?

```
new_bank = select(bank,c("approval","bool1","cont4","bool2","ages","cont6","cont3"))
head(new_bank)
```

```
## approval bool1 cont4 bool2 ages cont6 cont3
## 1      +      t      1      t  58      0  1.25
## 2      +      t      6      t  54     560  3.04
## 3      +      t      0      f  62     824  1.50
## 4      +      t      5      t  51       3  3.75
## 5      +      t      0      f  58       0  1.71
## 6      +      t      0      f  37       0  2.50
```

```
str(new_bank)
```

```
## 'data.frame': 666 obs. of 7 variables:
## $ approval: Factor w/ 2 levels "-","+": 2 2 2 2 2 2 2 2 2 2 ...
## $ bool1 : chr "t" "t" "t" "t" ...
## $ cont4 : int 1 6 0 5 0 0 0 0 0 0 ...
## $ bool2 : chr "t" "t" "f" "t" ...
## $ ages : int 58 54 62 51 58 37 47 67 61 62 ...
```

```
## $ cont6 : int 0 560 824 3 0 0 31285 1349 314 1442 ...
## $ cont3 : num 1.25 3.04 1.5 3.75 1.71 ...
## - attr(*, "na.action")= 'omit' Named int [1:24] 72 84 87 93 98 203 207 244 255 271 ...
## ..- attr(*, "names")= chr [1:24] "72" "84" "87" "93" ...
```

```
# split the data
index1 = createDataPartition(y=new_bank$approval, p=0.7, list=FALSE)
train_set = new_bank[index1,]
test_set = new_bank[-index1,]
```

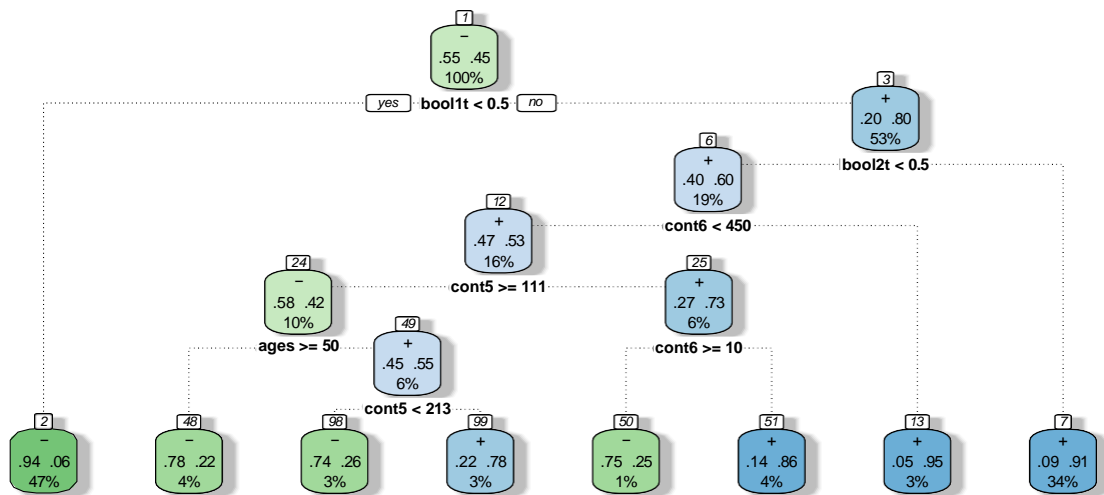
```
tree2 <- train(approval ~., data = train_set, method = "rpart1SE",
               trControl = train_control)
tree2
```

```
## CART
##
## 467 samples
## 6 predictor
## 2 classes: '-', '+'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 420, 421, 420, 420, 421, 421, ...
## Resampling results:
##
## Accuracy Kappa
## 0.8586031 0.7149616
```

we can observe that accuracy has increased after selecting relevant predictors (first 6).

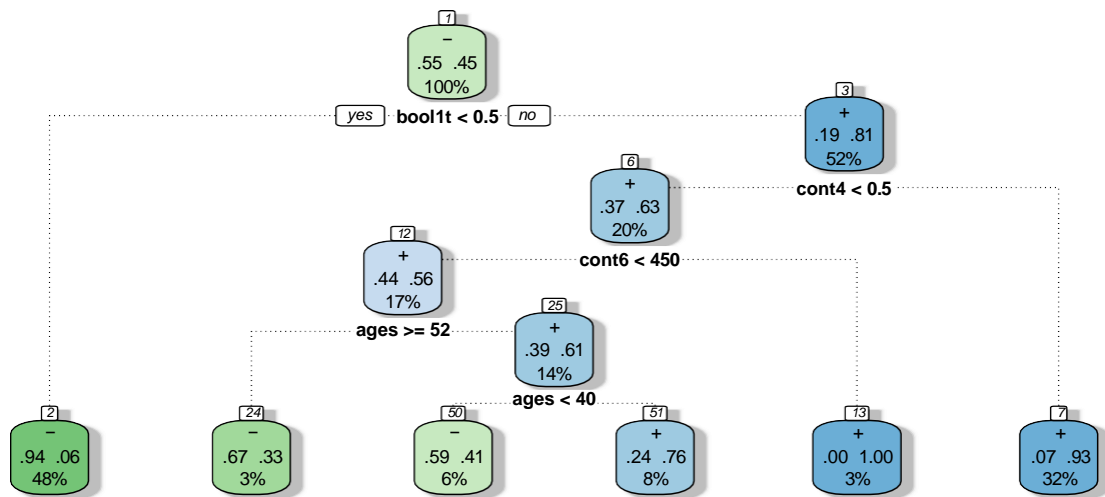
e. Visualize the trees from (a) and (d) and report if reducing the number of variables had an effect on the size of the tree?

```
fancyRpartPlot(tree$finalModel, caption= "tree1")
```

tree1

```
fancyRpartPlot(tree2$finalModel, caption= "tree2")
```



tree2

reducing the number of variables will reduce the graph.